

MA308 Mini Project
Report

Designing a shortest-path algorithm for large-scale graphs

Submitted by

Roll No.	Names of Students
----------	-------------------

I22MA023	Abhinav Kumar
I22MA038	Raj Kumar
I22MA062	Chandra Pratap

Under the guidance of
Dr. Sushil Kumar



Department of Mathematics
SARDAR VALLABHBHAI NATIONAL INSTITUTE OF TECHNOLOGY, SURAT
Even Semester 2024

Acknowledgement

We, the students of the 5-Year Integrated M.Sc. in Mathematics program at Sardar Vallabhbhai National Institute of Technology, Surat take this opportunity to express our profound gratitude to everyone who supported and guided us throughout this mini-project.

Our deepest thanks goes to **Dr. Sushil Kumar**, whose exceptional mentorship, guidance and encouragement have been pivotal to the success of this project. His ability to provide thoughtful insights and challenge us intellectually has greatly enriched our learning experience. Working under his supervision has been a rewarding and transformative experience that will benefit us throughout our academic and professional journeys.

We are equally grateful to **Dr. Anupam Shukla**, Director of SVNIT Surat, and **Dr. Jayesh M. Dhodiya**, Head of the Department of Mathematics, for their leadership and support in fostering an environment that prioritizes academic excellence and research innovation. We also extend our gratitude to all the faculty members, research scholars, and non-teaching staff of the department for their valuable assistance, constant encouragement, and willingness to help whenever needed.

This project was a collaborative effort, and we would like to acknowledge the contributions of one another, as well as those who supported us beyond the academic sphere:

I am deeply thankful to **Dr. Sushil Kumar** for his unwavering support, which has been instrumental in overcoming challenges during this project. I am grateful to my teammates for their cooperation and dedication, making this collaboration a fulfilling experience. Special thanks to my parents for their unconditional support and encouragement, as well as to all those who offered their assistance throughout this journey.

Abhinav Kumar

I extend my sincere gratitude to **Dr. Sushil Kumar**, whose guidance has been crucial in shaping my understanding of this subject. His mentorship has

made this project both insightful and rewarding. I am also immensely thankful to my teammates, **Abhinav Kumar** and **Chandra Pratap**, for their enthusiasm and teamwork. Finally, I wish to thank my family and friends for their encouragement, which has been a constant source of motivation during this project.

Raj Kumar

I express my heartfelt gratitude to **Dr. Sushil Kumar**, whose support and expertise have been invaluable in the successful completion of this project. I am thankful to my teammates, **Abhinav Kumar** and **Raj Kumar**, for their dedication, without which this project would not have been possible. Lastly, I am grateful to my family and friends, whose encouragement, patience, and unwavering belief in my abilities have been a source of strength throughout this journey.

Chandra Pratap

Department of Mathematics

SARDAR VALLABHBHAI NATIONAL INSTITUTE OF TECHNOLOGY, SURAT

Certificate

This is to certify that this is a bonafide record of the project presented by the students whose names are given below during the even semester of 2024 in partial fulfilment of the requirements of the degree of Integrated Masters of Science in Mathematics.

Roll No	Names of Students
---------	-------------------

I22MA023	Abhinav Kumar
I22MA038	Raj Kumar
I22MA062	Chandra Pratap

Dr. Sushil Kumar
(Project Guide)

Dr. Ramakanta Meher
(Course Coordinator)

Date: 20 January, 2025

Abstract

- Brief overview of the project.
- Key motivation, methodology, and outcomes.
- Highlights of the hybrid algorithm approach and its applications.

Contents

1	Problem Definition	1
1.1	Graph Representation	1
1.2	Objectives	1
2	Introduction	2
2.1	Background	2
2.2	Objective	2
2.3	Scope	2
3	Literature Review	3
3.1	Classical shortest path algorithms	3
3.1.1	Breadth-first Search	3
3.1.2	Bellman-ford Algorithm	4
3.1.3	Dijkstra’s Algorithm	5
3.2	Advanced shortest path algorithms	7
3.2.1	A* Search Algorithm	7
3.2.2	Bidirectional Search	8
3.3	Preprocessing techniques	8
3.3.1	Contraction Hierarchies	8
3.3.2	ALT Algorithm	8
3.3.3	Hub Labeling	8
3.4	Summary of Findings	8
3.5	Gaps in current approaches	8
4	Algorithm Design	10
4.1	Preprocessing Phase	10
4.2	Query Execution	10
4.3	Dynamic Updates	10
4.4	Customization	10

5	Implementation	11
5.1	Tools and Technologies	11
5.2	Code Structure	11
5.3	Dataset	11
6	Result and Analysis	12
6.1	Performance Metrics	12
6.2	Comparative Analysis	12
6.3	Sensitivity Analysis	12
7	Discussion	13
7.1	Strengths	13
7.2	Limitations	13
7.3	Further Improvements	13
8	Conclusion	14
	References	15
	Appendices	16
A		17
A.1	Proof of correctness for BFS	17
A.2	Proof of complexity for BFS	17
A.3	Proof of correctness for Bellman-Ford	18
A.4	Proof of complexity for Bellman-Ford	18
A.5	Proof of correctness for Dijkstra's	19
A.6	Proof of complexity for Dijkstra's	19
A.7	Proof of correctness for A* search	20

List of Figures

Chapter 1

Problem Definition

1.1 Graph Representation

- Define the graph $G=(V, E)$ where V is the set of vertices (nodes) and E is the set of edges.
- Edge weights represent costs (e.g., travel time, distance, tolls).

1.2 Objectives

- Minimize computation time for shortest path queries.
- Incorporate real-time edge weight updates.
- Allow customizable routing based on user-defined cost functions.

Chapter 2

Introduction

2.1 Background

- Importance of shortest path calculations in large-scale graphs.
- Challenges in real-world applications: scale, dynamics, and user preferences.

2.2 Objective

- To develop and analyze a hybrid shortest path algorithm combining preprocessing, efficient querying, and dynamic updates.

2.3 Scope

- Focus on road networks, real-time traffic updates, and customizable routing.

Chapter 3

Literature Review

In this chapter, we review the foundations and advanced algorithms for shortest path calculations, including preprocessing techniques. The section spans classical algorithms, advanced algorithms and recent advances in graph-optimization.

3.1 Classical shortest path algorithms

3.1.1 Breadth-first Search

Introduction

Breadth-first search is a graph traversal algorithms invented by Konrad Zuse in 1945, that can also be used to find the shortest path from a source vertex to a destination vertex in an unweighted graph.

Algorithm

1. Mark all vertices as unvisited.
2. Assign $distance[u] = \infty$ for all vertices except the source vertex s , where $distance[s] = 0$.
3. Use a queue to track vertices to explore. Start with the source vertex s .
4. Dequeue a vertex u .
5. For each neighbour v of u , If v is unvisited (i.e., $distance[v] = \infty$):
 - Set $distance[v] = distance[u] + 1$.

- Mark v as visited.
 - Enqueue v .
6. The algorithm ends when the queue is empty. Unreachable vertices retain $distance = \infty$.

This algorithm is mathematically predisposed to find the shortest path from a source vertex s to every other vertex in the graph (see **Appendix A.1** for a formal proof).

Complexity

When finding the shortest path between a pair of vertices in a graph, the worst-case time complexity for the BFS algorithm is $O(V)$ for queue operations + $O(E)$ for edge processing, netting a worst-case time complexity of $O(V + E)$ (see **Appendix A.2** for a formal proof).

The space complexity for BFS is $O(V)$ since we use a queue to store the vertices yet to be explored.

Pros and Cons

- The algorithm is simple and efficient for unweighted graphs.
- BFS works well for large, sparse graphs.
- BFS fails for shortest-path problems in weighted graphs, which are more useful when modelling real world scenarios.

3.1.2 Bellman-ford Algorithm

Introduction

The Bellman–Ford algorithm is a shortest-path algorithm that utilizes dynamic programming to compute shortest paths from a single source vertex to all of the other vertices in a weighted, directed graph. It was first published by Richard Bellman (1958) and Lester Ford Jr. (1956), hence its name.

Algorithm

1. Create an array $distance$ of size V to store the shortest path distances.
2. Assign $distance[u] = \infty$ for all vertices except the source vertex s , where $distance[s] = 0$.

3. Repeat $V - 1$ times:
 - For each edge $(u, v) \in E$, if $distance[u] + w(u, v) < distance[v]$ update $distance[v] = distance[u] + w(u, v)$.
4. Now to detect a negative cycle, for each edge $(u, v) \in E$, if $distance[u] + w(u, v) < distance[v]$, report that a negative-weight cycle exists.
5. If no negative-weight cycle is detected, Return the *distance* array as the shortest path distances.

Refer to **Appendix A.3** for a formal proof of correctness of this algorithm.

Complexity

When finding the shortest path from a source vertex to every other vertex in a graph, the worst-case time complexity for the Bellman-Ford algorithm is $O(V \cdot E)$ (see **Appendix A.4** for a formal proof).

The space complexity for Bellman-Ford is $O(V)$ since we use an array of size V to store all the shortest-path distances.

Pros and Cons

- Suitable for applications requiring negative weight handling, in which case it can detect the existence of a negative cycle.
- The Bellman-Ford algorithm is more complex than Dijkstra's algorithm.
- Much slower compared to Dijkstra's algorithm.

3.1.3 Dijkstra's Algorithm

Introduction

Dijkstra's algorithm is a greedy algorithm used to find the shortest paths from a single source vertex to all other vertices in a weighted graph with non-negative edge weights. It was conceived by computer scientist Edsger W. Dijkstra in 1956 and published three years later.

Algorithm

1. Create an array *distance* of size V to store the shortest path distances and a priority queue Q containing all vertices, prioritized by *distance*.
2. Assign $distance[u] = \infty$ for all vertices except the source vertex s , where $distance[s] = 0$.
3. While Q is not empty:
 - Extract the vertex u with the smallest distance from Q .
 - For each neighbor v of u , if $distance[u] + w(u, v) < distance[v]$: update $distance[v] = distance[u] + w(u, v)$ and the priority of v in Q .
4. The algorithm ends when Q is empty. The distance array contains the shortest path distances from s to all other vertices.

Refer to **Appendix A.5** for a formal proof of correctness of this algorithm.

Complexity

When finding the shortest path from a source vertex to every other vertex in a graph, the worst-case time complexity for the Dijkstra's algorithm is $O((V + E) \log V)$ using a binary heap or $O(V^2 + E)$ using a Fibonacci heap. (see **Appendix A.6** for a formal proof).

The space complexity for Dijkstra's is $O(V)$ since we use an array of size V to store all the shortest-path distances.

Pros and Cons

- Can cover a large area of a graph, which is useful when there are multiple target nodes.
- Can't calculate the shortest paths correctly if the graph has negative weights.
- Has linearithmic complexity when implemented using a priority queue.

3.2 Advanced shortest path algorithms

3.2.1 A* Search Algorithm

Introduction

A* search is a heuristic-based algorithm used to find the shortest path from a start node to a goal node in a weighted graph. It combines the strengths of Dijkstra's algorithm (guaranteed shortest path) and greedy best-first search (efficient exploration using heuristics). It was first published by Peter Hart, Nils Nilsson, and Bertram Raphael at Stanford Research Institute in 1968.

Algorithm

1. Create a priority queue Q to store nodes to explore, prioritized by $f(v) = g(v) + h(v)$, where
 - $g(v)$: Cost of the shortest path from s to v found so far.
 - $h(v)$: Heuristic estimate of the cost from v to t .
2. Set $g(s) = 0$ and $f(s) = h(s)$.
3. Insert s into Q .
4. Create a set *visited* to track visited nodes
5. While Q is not empty:
 - (a) Extract the node u with the smallest $f(u)$ from Q .
 - (b) If $u = t$, return the path from s to t .
 - (c) Mark u as visited.
 - (d) For each neighbor v of u , if v is not visited:
 - Compute $g_{tentative} = g(u) + w(u, v)$.
 - If $g_{tentative} < g(v)$ or v is not in Q :
 - Update $g(v) = g_{tentative}$.
 - Update $f(v) = g(v) + h(v)$.
 - Insert v into Q (or update its priority if already in Q).
6. If Q becomes empty and the goal t has not been reached, no path exists.

A* search is correct if the heuristic $h(v)$ is admissible (never overestimates the true cost to the goal) and consistent (satisfies the triangle inequality: $h(u) \leq w(u, v) + h(v)$ for all edges (u, v)). For a formal proof of correctness, refer to **Appendix A.7**.

Complexity

Since A* Search is basically an 'informed' version of Dijkstra's algorithm, the space complexity for A* search is the same as for Dijkstra's, which is $O(V)$. The time complexity, however, depends on the heuristic function and is equal to Dijkstra's when the heuristic $h(v) = 0$.

Pros and Cons

- Compared to uninformed search algorithms, A* explores significantly fewer nodes leading to faster search times.
- By maintaining a priority queue, A* only needs to store a limited number of nodes in memory, making it suitable for large search spaces.
- Performance heavily depends on the quality of the heuristic function. Thus, A* search is not ideal when a good heuristic cannot be easily defined or when heuristic calculations are complicated.

3.2.2 Bidirectional Search

Introduction

Algorithm

Complexity

Pros and Cons

3.3 Preprocessing techniques

3.3.1 Contraction Hierarchies

3.3.2 ALT Algorithm

3.3.3 Hub Labeling

Refer figure ??.

3.4 Summary of Findings

3.5 Gaps in current approaches

- Limited scalability in dynamic graphs.

- Lack of flexibility for user preferences.
- High preprocessing overhead in large graphs.

Chapter 4

Algorithm Design

4.1 Preprocessing Phase

- Explain the choice of preprocessing techniques (e.g., Contraction Hierarchies, ALT).
- Describe how the graph is simplified or augmented with shortcuts/landmarks.
- Discuss trade-offs in preprocessing time and memory usage.

4.2 Query Execution

- Detail the hybrid algorithm combining preprocessing results with A* or bidirectional search.
- Include pseudocode for the query phase.

4.3 Dynamic Updates

- Describe methods for updating edge weights and recalculating shortest paths efficiently.
- Outline incremental update mechanisms.

4.4 Customization

- Define the composite cost function and how it integrates with the algorithm.

Chapter 5

Implementation

5.1 Tools and Technologies

- Programming language and libraries/frameworks used (e.g., Python, NetworkX, CUDA for parallelism).
- Hardware setup for experiments.

5.2 Code Structure

- Modular design: preprocessing, query execution, dynamic updates, and customization.

5.3 Dataset

- Description of datasets used (e.g., OpenStreetMap, SNAP datasets).
- Data preprocessing steps: parsing, cleaning, and formatting.

Chapter 6

Result and Analysis

6.1 Performance Metrics

- Preprocessing time and memory usage.
- Query execution time (average, worst-case).
- Accuracy of paths (for heuristic methods).
- Scalability with graph size and density.

6.2 Comparative Analysis

- Benchmark hybrid algorithm against standalone algorithms (e.g., plain Dijkstra's, A*).
- Use tables and plots to illustrate improvements.

6.3 Sensitivity Analysis

- Impact of graph properties (e.g., number of nodes, edge density).
- Effect of dynamic updates on performance.

Chapter 7

Discussion

7.1 Strengths

- Highlight significant improvements in speed, accuracy, or flexibility.
- Discuss adaptability to various real-world scenarios.

7.2 Limitations

- Discuss preprocessing overhead or constraints in memory usage.
- Identify scenarios where the hybrid approach might underperform.

7.3 Further Improvements

- Enhancing scalability for distributed systems.
- Incorporating machine learning to predict edge weights dynamically.
- Extending to multimodal routing or 3D graphs.

Chapter 8

Conclusion

- Recap the problem, approach, and key findings.
- Reiterate the significance of combining multiple algorithms for shortest path calculations.
- Highlight practical implications and potential impact.

References

- [1] Cite all academic papers, libraries, datasets, and tools used. `<urlhere>`

Appendices

Appendix A

A.1 Proof of correctness for BFS

We'll prove the correctness of BFS using mathematical induction.

- *Inductive hypothesis:* For all nodes at distance k from the source, BFS correctly computes $distance[v] = k$.
- *Base case:* The source node s has $distance[s] = 0$.
- *Induction step:* Assume the hypothesis is true for nodes at a distance k from s . Then their neighbours (nodes at distance $k + 1$) are enqueued and assigned $distance = k + 1$ before any nodes at $distance > k + 1$ are processed.
- *Conclusion:* BFS computes the shortest possible path for all reachable nodes.

A.2 Proof of complexity for BFS

Let us assume a graph $G(V, E)$ with V vertices and E edges.

- Mark all V vertices as unvisited. This takes $O(V)$ time.
- Each vertex enters the queue once (when discovered) and exits the queue once. Enqueue and dequeue operations are $O(1)$, so processing all vertices takes $O(V)$ time.
- For each dequeued vertex u , iterate through its adjacency list to check all edges (u, v) .
- In a directed graph, each edge (u, v) is processed once. In an undirected graph, each edge (u, v) is stored twice (once for u and once for v), but each is still processed once during BFS.

- Summing over all vertices, the total edge-processing time is $O(E)$.

Thus, the overall time complexity is $O(V + E)$.

A.3 Proof of correctness for Bellman-Ford

We'll prove the correctness of Bellman-Ford algorithm using mathematical induction.

- *Inductive hypothesis:* After k iterations, $distance[v]$ is the length of the shortest path from s to v using at most k edges.
- *Base case:* After 0 iterations, $distance[s] = 0$ (correct), and $distance[v] = \infty$ for all $v \neq s$ (no paths have been explored yet).
- *Induction step:* Consider the $(k + 1)^{th}$ iteration. For each edge (u, v) , if $distance[u] + w(u, v) < distance[v]$, then $distance[v]$ is updated to $distance[u] + w(u, v)$. This ensures that after $k + 1$ iterations, $distance[v]$ is the length of the shortest path using at most $k + 1$ edges.
- *Conclusion:* After $V - 1$ iterations, all shortest paths with at most $V - 1$ edges have been found. Since a shortest path in a graph with V vertices cannot have more than $V - 1$ edges, the algorithm is correct.
- *Negative cycle detection:* After $V - 1$ iterations, if any $distance[v]$ can still be improved (i.e. $distance[u] + w(u, v) < distance[v]$ for some edge (u, v)), then the graph contains a negative-weight cycle reachable from s .

A.4 Proof of complexity for Bellman-Ford

Let us assume a graph $G(V, E)$ with V vertices and E edges.

- Set $distance[s] = 0$ and $distance[v] = \infty$ for all $v \neq s$. This takes $O(V)$ time.
- Relax all E edges, repeated $V - 1$ times. Each relaxation takes $O(1)$ time. This takes a total time of $O(V \cdot E)$.
- For negative cycle detection, relax all the edges once more. This takes $O(E)$ time.

The dominant term is the relaxation step, which takes $O(V \cdot E)$ time, hence the overall complexity of the algorithm is $O(V \cdot E)$.

A.5 Proof of correctness for Dijkstra's

We'll prove the correctness of Dijkstra's algorithm using mathematical induction.

- *Inductive hypothesis:* After k vertices are extracted from Q , their distance values are the correct shortest path distances from s .
- *Base case:* Initially, $distance[s] = 0$ (correct), and $distance[v] = \infty$ for all $v \neq s$ (no paths have been explored yet).
- *Induction step:* Let u be the $(k + 1)^{th}$ vertex extracted from Q . Suppose there exists a shorter path to u not using the extracted vertices. This path must leave the set of extracted vertices at some edge (x, y) , but since $w(x, y) \geq 0$, this would imply $distance[y] < distance[u]$, contradicting u 's extraction.
- *Conclusion:* After all vertices are processed, the *distance* array contains the correct shortest path distances.

A.6 Proof of complexity for Dijkstra's

Let us assume a graph $G(V, E)$ with V vertices and E edges. In a priority-queue based implementation of the algorithm,

- Each vertex is extracted once ($V \times \text{Extract-Min}$) and each edge is relaxed once ($E \times \text{Decrease-Key}$).
- Extract-Min and Decrease-Key take $O(\log V)$ time in a binary heap.
- Extract-Min and Decrease-Key take $O(\log V)$ and $O(1)$ time respectively in a fibonacci heap.
- For a binary heap, $V \times \text{Extract-Min}$ takes $O(V \log V)$ time and $E \times \text{Decrease-Key}$ takes $O(E \log V)$ time \rightarrow a total complexity of $O((V + E) \log V)$
- For a fibonacci heap, $V \times \text{Extract-Min}$ takes $O(V \log V)$ time and $E \times \text{Decrease-Key}$ takes $O(E)$ time \rightarrow a total complexity of $O(V \log V + E)$.

A.7 Proof of correctness for A* search

We'll prove the correctness of A* search algorithm using mathematical induction. Let us define the following:

$f(s)$: Estimated total cost of the path from the start node to the goal node, passing through the current node.

$g(s)$: Cost of the shortest path from the start node to the current node.

$h(s)$: Heuristic estimate of the cost from the current node to the goal node.

- *Inductive hypothesis*: At each step, the node u with the smallest $f(u)$ is the one with the smallest estimated total cost to the goal.
- *Base case*: Initially, $g(s) = 0$ and $f(s) = h(s)$. The start node s is correctly prioritized.
- *Induction step*:
 - When u is extracted, its $g(u)$ is the true shortest path cost from s to u (due to admissibility and consistency).
 - For each neighbor v , $f(v) = g(v) + h(v)$ is updated to reflect the best-known path to v .
 - The algorithm continues to explore nodes in order of increasing $f(v)$, ensuring the shortest path is found.
- *Conclusion*: If the goal t is reached, $g(t)$ is the true shortest path cost and If Q becomes empty, no path exists.